

Task 5: Image Stitching & RANSAC

Name: Laurent G  rin Degree: MSc ID: 20505439

Problem 1: Separable Filter

(a) What is a separability of the Gaussian kernel? When do we use this separability and why does this characteristic reduce the computing cost?

The separability of the Gaussian kernel means that the kernel can be represented as the multiplication of two vectors. For instance, a 3x3 Gaussian kernel, we solve the system of equations:

$$G = \begin{bmatrix} 0.0113 & 0.0838 & 0.0113 \\ 0.0838 & 0.6193 & 0.0838 \\ 0.0113 & 0.0838 & 0.0113 \end{bmatrix} = \begin{bmatrix} a \\ b \\ c \end{bmatrix} \begin{bmatrix} d & e & f \end{bmatrix} = \begin{bmatrix} ad & ae & af \\ bd & be & bf \\ cd & ce & cf \end{bmatrix}$$

First extracting the middle column vectors from both matrices, we get three relationships:

$$\begin{bmatrix} 0.0838 \\ 0.6193 \\ 0.0838 \end{bmatrix} = \begin{bmatrix} ae \\ be \\ ce \end{bmatrix}$$
$$\therefore \begin{aligned} a &= 0.0838/e, \\ b &= 0.6193/e, \\ c &= 0.0838/e \end{aligned}$$

Similarly with the middle row vectors:

$$\begin{bmatrix} 0.0838 & 0.6193 & 0.0838 \end{bmatrix} = \begin{bmatrix} bd & be & bf \end{bmatrix}$$
$$\therefore \begin{aligned} d &= 0.0838/b, \\ f &= 0.0838/b \end{aligned}$$

We can substitute b :

$$\begin{aligned} d &= 0.1353e, \\ f &= 0.1353e \end{aligned}$$

We are thus left with this:

$$G = \begin{bmatrix} 0.0838/e \\ 0.6193/e \\ 0.0838/e \end{bmatrix} \begin{bmatrix} 0.1353e & e & 0.1353e \end{bmatrix}$$

Removing e from both vectors, it is clear that e cancels out. Thus we obtain our final separated kernel:

$$G = \begin{bmatrix} 0.0838 \\ 0.6193 \\ 0.0838 \end{bmatrix} \begin{bmatrix} 0.1353 & 1 & 0.1353 \end{bmatrix}$$

This separability feature becomes quite useful when performing convolution to blur an image. We can define an original image I and a blurred image I_b , the result of convolution with the Gaussian kernel:

$$I_b = G * I$$

Let us now separate G , where G_1 is the vertical vector and G_2 is the horizontal vector:

$$I_b = (G_1 G_2) * I$$

It turns out that $G_1 G_2$ is the same as finding their 2D convolution, i.e. $G_1 * G_2$ (http://www.songho.ca/dsp/convolution/convolution2d_separable.html). We can also use the property of commutativity of convolution, and rewrite:

$$I_b = (G_1 * G_2) * I = G_1 * G_2 * I = G_1 * (G_2 * I)$$

This brings significant computational savings. If I is of size $m \times n$, and G , of size $\ell \times \ell$, regular 2D convolution would take $\ell^2 mn$ multiplications, and $(\ell^2 - 1)mn$ additions (since for each pixel we multiply each neighbours reached by G with their respective weight and then sum them). This means that larger Gaussian kernels will very rapidly increase the computation cost.

However, if we first convolve with G_2 , which has size ℓ , convolution takes ℓmn multiplications and $(\ell - 1)mn$ additions. Then if we convolve this result with G_1 , we again need ℓmn multiplications and $(\ell - 1)mn$ additions. Thus we get a total of $2\ell mn$ multiplications and $2(\ell - 1)mn$ additions.

In Big-O notation, the first method, naive 2D convolution, has complexity $O(\ell^2 mn)$, while the second method, using the separated kernel, has complexity $O(\ell mn)$. This is huge! Now the computational time only increases linearly with the size of the kernel instead of quadratically.

(b) What is the Laplacian of Gaussian (LoG)? When do we use LoG?

The Laplacian of Gaussian is a kernel which allows two operations at the same time. Essentially, we take the 2nd derivative of our image (Laplacian), and blur it out (Gaussian). We blur it out because 2nd derivatives are extremely sensitive to noise. The LoG allows combining the two operations at once and thus saves computational requirements.

The LoG becomes useful for corner and blob detection - since it has a round shape and strongly defined edges, local extrema after the convolution will indicate where strong blobs were detected.

(c) Is LoG separable?

The LoG is not separable. As an example, if we look at a 5x5 LoG kernel:

$$LoG = \begin{bmatrix} 0 & 0 & 1 & 0 & 0 \\ 0 & 1 & 2 & 1 & 0 \\ 1 & 2 & -16 & 2 & 1 \\ 0 & 1 & 2 & 1 & 0 \\ 0 & 0 & 1 & 0 & 0 \end{bmatrix}$$

We can see that the rows are clearly not linear combinations of each other, as they are in the Gaussian kernel. ie, there is no way to multiply a row-vector by scalars a and b that will give rows 1 and 2 in the matrix.

(d) What is a difference of the Gaussian (DoG)? When do we use it and what is DoG advantageous compared to LoG?

The difference of Gaussian is very similar to LoG, with a slightly less "pointy" blob. Essentially, we take the Gaussian blur of an image with a certain σ , and then we change σ , take another Gaussian blur, and find the difference between the two images. The advantage of DoG is that we can use the Gaussian kernel's separability to accelerate processing, and also that since we are taking the Gaussian with two different σ values, we can quickly look at different blob sizes in a "chain" manner, further reducing the computational time compared to the LoG.

Problem 2: Least Squares

(a) Explain the approach 1 and approach 2 for least squares line fitting in your words. Please refer to the course slide and tutorials.

Both approaches are mathematically the same, the only difference being that approach 2 represents the least-square fitting method in matrix form. In both cases, we are minimizing the error between the fitted model and the data. First, we define an error term:

$$E = \sum_{i=1}^n (y_i - mx_i - b)^2$$

or in matrix form (approach 2):

$$E = ||Y - XB||^2$$

This error term is the difference between the measured y values and their predicted value by the model. Squaring this difference, we eliminate negative values and thus avoid the case where negative error terms cancel positive error terms. We also put more emphasis on values which have larger error terms, ie which are further away from our fitted line.

In approach 2, this is defined as the Euclidian norm, i.e. the "distance" from the origin in n dimensions. Since the distance is calculated as such:

$$||Y - XB|| = \sqrt{(y_1 - (mx_1 + b))^2 + (y_2 - (mx_2 + b))^2 + \dots}$$

We square this to remove the square root, and thus we can see how this is exactly the same as the error term in approach 1.

After defining our error term, we take the partial derivatives with respect to m and b , or in approach 2, B . Setting these equal to zero and solving for m and b or B , we will find the model which minimizes our error term:

$$m = \frac{\sum x_i y_i - (1/n) \sum x_i \sum y_i}{\sum x_i^2 - (1/n)(\sum x_i)^2}$$

$$b = \frac{\sum y_i \sum x_i^2 - \sum x_i \sum x_i y_i}{n \sum x_i^2 - (\sum x_i)^2}$$

$$B = (X^T X)^{-1} X^T Y$$

A response y_i is measured from a linear system when an input is x_i . The measurement data is provided (`prob2_data1.mat`). A model of your system can be approximated as $y_i = mx_i + b$. Please find m and b using the following methods.

(b) Least squares (approach 1)

(c) Least squares (approach 2)

We find with both methods that m is 1.6540 and b is 2.3197.

```

1  load('task_files/prob2_data1.mat')
2
3  %approach 1:
4  n = length(x);
5  SX = sum(x);
6  SY = sum(y);
7  SXY = sum(x .* y);
8  SX2 = sum(x.^2);
9  S2X = SX^2;
10
11 m = (SXY - 1/n * SX * SY)/(SX2 - 1/n * S2X)
12 b = (SY * SX2 - SX * SXY)/(n * SX2 - S2X)
13
14 %approach 2:
15 X = [x; ones(1,n)]';
16 B = (X' * X) \ X' * y'
```

Here is another measurement data (`prob2_data2.mat`). Please find m and b using the following methods.

(d) Least squares (either approach 1 or approach 2)


We find $m = 2.3348$ and $b = 5.6458$.

```

1  load('task_files/prob2_data1.mat')
2  X = [x; ones(1,length(x))];
3  B = (X' * X) \ X' * y'
```

(e) Use of RANSAC. You need to have your own RANSAC implementation without using existing functions in MATLAB (Do not use `ransac` or any other relevant functions)


We find $m = 2.3074$ and $b = 5.5433$. Comparison of the two fitted lines shows this:



Problem 3: Fitting using RANSAC

(a) Fit an ellipse to the given data (`prob3_ellipse.mat`) using RANSAC

The ellipse was successfully fitted using RANSAC as shown here.



```

1 %Fit an ellipse
2 %Using general quadratic curve  $ax^2 + bxy + cy^2 + dx + ey + f = 0$ 
3
4 load('task_files/prob3_ellipse.mat')
5 plot(x,y);
6
7 n_tries = 10000;
8 thres = 0.1;
9
10 combinations = zeros(n_tries,7); %a, b, c, d, e, f, count
11
12 for i = 1:n_tries
13     [xi, idx] = datasample(x,5);
14     yi = y(idx);
15
16     bigMatrix = zeros(5,6);
17     for j = 1:5
18         bigMatrix(j,:) = [xi(j)^2, xi(j) * yi(j), yi(j)^2, xi(j), yi(j), 1];
19     end
20     params = null(bigMatrix); %a, b, c, d, e, f
21     if size(params,2) > 1
22         params = params(:,1);
23     end
24     %we have an implicit equation which makes it difficult to find the vertical
    distance. Instead we use the

```


```

25 %value obtained by putting x and y into the ellipse equation, which should be
    zero.
26 count = sum( abs(params(1)*x.^2 + params(2)*x.*y + params(3)*y.^2 + params(4)*x +
    params(5)*y + params(6)) <= thres);
27 combinations(i,:)=[params',count];
28 end
29
30 [maxcount, maxidx] = max(combinations(:,7));
31 a = combinations(maxidx,1);
32 b = combinations(maxidx,2);
33 c = combinations(maxidx,3);
34 d = combinations(maxidx,4);
35 e = combinations(maxidx,5);
36 f = combinations(maxidx,6);
37 hold on
38 fimplicit(@(x,y) a*x.^2 + b*x.*y + c*y.^2 + d*x + e*y + f,'r--','Linewidth',2);

```

(b) Fit a fourth degree polynomials to the given data (prob3_polynomial.mat) using RANSAC

The polynomial is successfully fitted as shown:



```

1 %Fit 4th order polynomial
2 %y = ax^4 + bx^3 + cx^2 + dx + e
3
4 load('task_files/prob3_polynomial.mat')

```

```

5 plot(x,y);
6
7 n_tries = 10000;
8 thres = 0.1;
9
10 combinations = zeros(n_tries,6); %a, b, c, d, e, count
11
12 for i = 1:n_tries
13     [xi, idx] = datasample(x,5);
14     yi = y(idx);
15
16     bigMatrix = zeros(5,5);
17     for j = 1:5
18         bigMatrix(j,:) = [xi(j)^4, xi(j)^3, xi(j)^2, xi(j), 1];
19     end
20     params = bigMatrix\yi'; %a, b, c, d, e
21
22     count = sum( abs(y - (params(1)*x.^4 + params(2)*x.^3 + params(3)*x.^2 +
23     params(4)*x + params(5))) <= thres);
24     combinations(i,:)=[params',count];
25 end
26
27 [maxcount, maxidx] = max(combinations(:,6));
28 a = combinations(maxidx,1);
29 b = combinations(maxidx,2);
30 c = combinations(maxidx,3);
31 d = combinations(maxidx,4);
32 e = combinations(maxidx,5);
33
34 hold on
35 y_fit = a * x.^4 + b*x.^3 + c*x.^2 + d*x + e;
36 plot(x,y_fit, 'r--','Linewidth',2)

```

Problem 4: Improved 3D Planar Measurement Tool

(a) Build your own measurement tool and evaluate your measurement using the images provided (see the folder of `prob4_img`). You may need to estimate homography based on SIFT feature matching. The exact size of the booklet is 24 cm x 31.5 cm and use `cover.jpg` to solve this problem.

The table below shows each image and the measurement taken to be from the 0 mark to the 10 cm mark on the ruler in the pictures. Note that inaccuracies also come from selection of the points on the picture.



Image	Measurement (should be 10 cm)
	10.11
	9.64





Image	Measurement (should be 10 cm)
	9.81
	9.50

Image	Measurement (should be 10 cm)
	10.11
	10.15

```

1 cover = imread('task_files/cover.jpg');
2 cover = imresize(cover, [1000, NaN]); %Resize to have 1000 rows
3 coverBW = single(rgb2gray(cover));
4
5
6 %fileName = char(strcat(num2str(numpictures,'task_files/%03.f'),".JPG"));
7 fileName = 'task_files/prob4_img/006.JPG';
8 img = imread(fileName);

```

```

9  img = imresize(img, [1000, NaN]); %Resize to have 1000 rows
10 imgBW = single(rgb2gray(img));
11
12 %detect cover features
13 [f_cover,d_cover] = vl_sift(coverBW);
14 if(0)
15     imshow(cover);
16     hold on
17     h1 = vl_plotframe(f_cover) ;
18     h2 = vl_plotframe(f_cover) ;
19     set(h1,'color','k','linewidth',3) ;
20     set(h2,'color','y','linewidth',2) ;
21 end
22
23 %detect image features
24 [f_img,d_img] = vl_sift(imgBW);
25 if(0)
26     %show these image features
27     figure
28     imshow(img);
29     hold on
30     h1 = vl_plotframe(f_img) ;
31     h2 = vl_plotframe(f_img) ;
32     set(h1,'color','k','linewidth',3) ;
33     set(h2,'color','y','linewidth',2) ;
34 end
35 %Match features
36 [matches, scores] = vl_ubcmatch(d_cover, d_img) ;
37
38 numMatches = length(scores);
39
40 %Get matching X and Y values
41 Pts_cover = [f_cover(1:2,matches(1,:));ones(1,numMatches)]; % add ones at the end
    to get homogeneous coordinates
42 Pts_img = [f_img(1:2,matches(2,:));ones(1,numMatches)];
43
44
45
46 %Now we estimate the homography
47 n_tries = 1000;
48 thres = 2; %pixels: distance to be considered outlier
49 %tform =
    estimateGeometricTransform(Pts_img',Pts_cover','projective','MaxNumTrials',n_tries,'M
    axDistance',5);
50
51 score = zeros(n_tries,1);
52 H = zeros(3,3,n_tries);
53 for i = 1:n_tries
54     x = zeros(1,4);
55     y = zeros(1,4);
56     xprime = zeros(1,4);
57     yprime = zeros(1,4);
58     %Find 4 points at random to estimate homography

```

```

59     for j = 1:4
60         index = randi(numMatches);
61         x(j) = Pts_cover(1,index);
62         y(j) = Pts_cover(2,index);
63         xprime(j) = Pts_img(1,index);
64         yprime(j) = Pts_img(2,index);
65     end
66
67     %Find homography matrix for these 4 points
68     bigMatrix=zeros(8,9);
69     for j=1:2:7
70         k = (j+1)/2;
71         bigMatrix(j:j+1,:) = [x(k) y(k) 1 0 0 0 -x(k)*xprime(k) -xprime(k)*y(k) -
xprime(k);...
72             0 0 0 x(k) y(k) 1 -x(k)*yprime(k) -yprime(k)*y(k) -yprime(k)];
73     end
74     h = null(bigMatrix);
75     h = h(:,1);
76     H(:, :, i) = reshape(h,3,3);
77
78     Pts_coverH = (Pts_cover' * H(:, :, i))'; %Transform cover points to img homography
79     dx = Pts_coverH(1,:)./Pts_coverH(3,:)-Pts_img(1,:);
80     dy = Pts_coverH(2,:)./Pts_coverH(3,:)-Pts_img(2,:);
81     score(i) = sum(dx.^2+dy.^2 <= thres^2);
82 end
83
84 [best, bestidx] = max(score);
85 H = H(:, :, bestidx);
86
87 %Rectify and show image
88 if(0)
89     imgH = imwarp(img,projective2d(inv(H)));
90     figure
91     imshow(imgH)
92     hold on
93 end
94
95 % %Bring cover to image plane
96 % coverH = imwarp(cover,projective2d(H));
97 % figure
98 % imshow(coverH)
99
100 %Now we know that the cover has dimensions 24 cm x 31.5 cm. We can use H
101 %to find the scaling factor to be used for our measurements.
102
103 dimX = 24;      %cm
104 dimY = 31.5;    %cm
105
106 figure
107 imshow(img)
108 hold on
109
110 %Find scale of cover picture

```



```

111 scX = 24/size(cover,2);    %cm/pixel
112 scY = 31.5/size(cover,1); %cm/pixel
113
114 sc = mean([scX,scY]);
115
116 %First line set
117 [xline1, yline1] = getline(); %Get first polyline - select vertices and hit enter
    once all are chosen
118 plot(xline1,yline1,'b','Linewidth',2)    %Show chosen lines
119
120 ptsLine1 = [xline1';yline1';ones(1,length(xline1))];
121
122 ptsLineRect1 = H' \ ptsLine1;
123 ptsLineRect1(1:2,:) = ptsLineRect1(1:2,:)./ptsLineRect1(3,:);
124
125 dist = zeros(length(xline1)-1,1);
126 for i = 1:(length(xline1)-1)
127     dist(i) = sqrt((ptsLineRect1(1,i) - (ptsLineRect1(1,i+1)))^2 + (ptsLineRect1(2,i)
    - (ptsLineRect1(2,i+1)))^2) * sc;
128     formatSpec = 'Length of line %1u is %2.3f cm\n';
129     fprintf(formatSpec,i,dist(i))
130 end

```

(b) Prepare your own calibration paper and take photos similar to the ones in (a). Then, evaluate your tool.

This was done using a fridge magnet from UW's Science Society... Dimensions were measured as **8.91 x 5.08 cm**. The table below shows a few pictures and measurements on the ruler, again taken at 10 cm. In general, the measured dimensions are not as good as in the previous picture. It's probably because the magnet is much smaller, thus we get a higher relative error. Additionally, this doesn't correct for lens distortion - this also adds additional error, particularly in the first picture where we can see some lines aren't parallel.




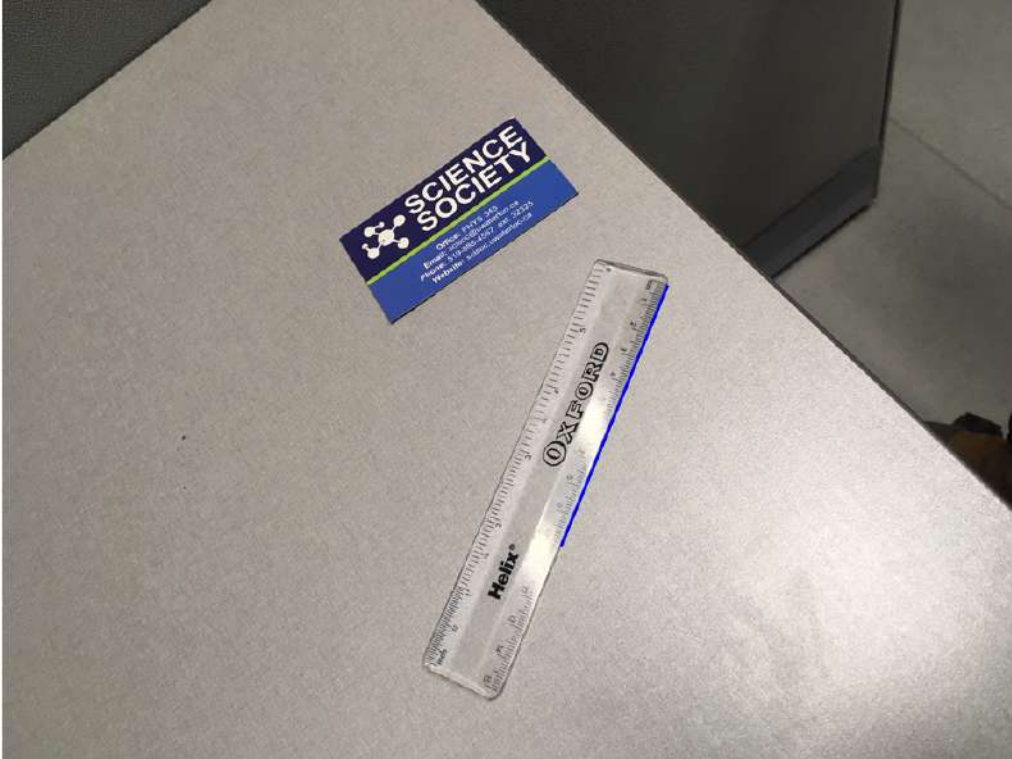
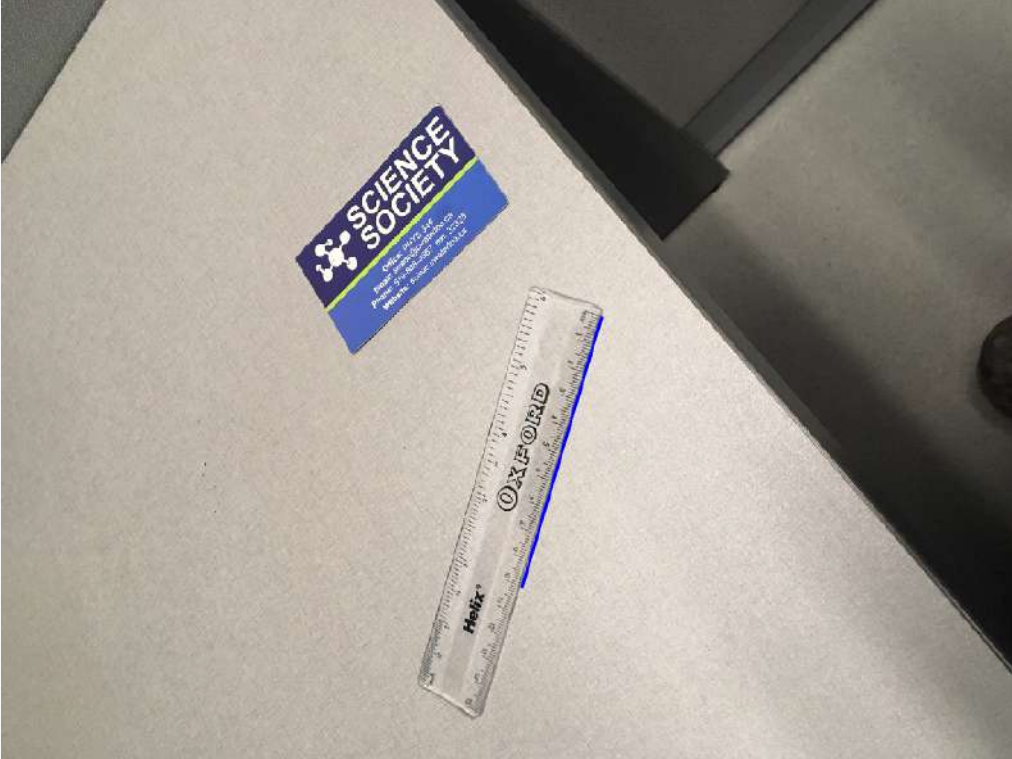
Image	Measurement (should be 10 cm)
	11.912
	9.419





Image	Measurement (should be 10 cm)
	9.978

(c) Compare measurements using this new tool and the one developed in Task 3.

In general, the measurements were at similar levels of accuracy. This is mostly because there is large imprecision in the selection of points with the mouse cursor. The additional distortion from camera parameters and the lens adds additional error as well. Also, since in this case we are doing feature-based matching for points, we can't exploit the full size of our "marker". In Task 3, we were using all 4 corners of the sheet, which made for larger distances which could "swallow" some of the error, whereas since features may be very close to each other, there is potential for small errors to be magnified once the homography is applied to the whole image.

Problem 5: Book Classification using SIFT

You are going to categorize books on images. Here are the input images and expected outcomes.

Test image	Result
	
	


Your code needs to automatically compute the outlines (boundary) of each book and its identity (book name). There are 17 images (see the folder of `prob5_img`) and in each image, four books are placed on a desk. Your code should not fail to identify a book or estimate its boundary more than 5 books among all books ($85 = 17 \text{ images} \times 5 \text{ books}$). Note that you should not do hard coding as well.

This was done using SIFT. Images, with all books identified, are shown in the table below:

Original



Classified



Original



Classified



Original



Classified



Original



Classified



Original



Classified



```

1  display = 0;
2  covers{1} = imread('covers_prob5/deep_goodfellow.jpg');
3  covers{2} = imread('covers_prob5/handbook_bovik.jpg');
4  covers{3} = imread('covers_prob5/pattern_bishop.jpg');
5  covers{4} = imread('covers_prob5/modern_mikhail.jpg');
6
7  names = {'deep-goodfellow', 'handbook-bovik', 'pattern-bishop', 'modern-mikhail'};
8
9  %detect cover features
10 for i = 1:length(covers)
11     coversBW{i} = single(rgb2gray(covers{i}));
12     [f_covers{i},d_covers{i}] = vl_sift(coversBW{i});
13     if(display)
14         figure
15         imshow(covers{i});
16         hold on
17         h1 = vl_plotframe(f_covers{i}) ;
18         h2 = vl_plotframe(d_covers{i}) ;
19         set(h1,'color','k','linewidth',3) ;
20         set(h2,'color','y','linewidth',2) ;
21     end
22 end
23
24 for nn = 1:17
25     fileName = char(strcat(num2str(nn,'task_files/prob5_img/%03.f'),'.JPG'));
26     img = imread(fileName);
27     img = imresize(img, [1000, NaN]); %Resize to have 1000 rows
28     imgBW = single(rgb2gray(img));
29
30 %detect image features
31 [f_img,d_img] = vl_sift(imgBW);
32 if(display)
33     %show these image features
34     figure
35     imshow(img);
36     hold on
37     h1 = vl_plotframe(f_img) ;
38     h2 = vl_plotframe(d_img) ;

```

```

39     set(h1,'color','k','linewidth',3) ;
40     set(h2,'color','y','linewidth',2) ;
41 end
42 %Match features
43 for n = 1:length(covers)
44     [matches, scores] = vl_ubcmatch(d_covers{n}, d_img);
45     numMatches = length(scores);
46     %Get matching X and Y values
47     Pts_cover = [f_covers{n}(1:2,matches(1,:));ones(1,numMatches)]; % add ones at
the end to get homogeneous coordinates
48     Pts_img = [f_img(1:2,matches(2,:));ones(1,numMatches)];
49
50     %Now we estimate the homography
51     n_tries = 1000;
52     thres = 3; %pixels: distance to be considered outlier
53
54     score = zeros(n_tries,1);
55     H = zeros(3,3,n_tries);
56     for i = 1:n_tries
57         x = zeros(1,4);
58         y = zeros(1,4);
59         xprime = zeros(1,4);
60         yprime = zeros(1,4);
61         %Find 4 points at random to estimate homography
62         for j = 1:4
63             index = randi(numMatches);
64             x(j) = Pts_cover(1,index);
65             y(j) = Pts_cover(2,index);
66             xprime(j) = Pts_img(1,index);
67             yprime(j) = Pts_img(2,index);
68         end
69
70         %Find homography matrix for these 4 points
71         bigMatrix=zeros(8,9);
72         for j=1:2:7
73             k = (j+1)/2;
74             bigMatrix(j:j+1,:) = [x(k) y(k) 1 0 0 0 -x(k)*xprime(k) -xprime(k)*y(k) -
xprime(k);...
75                 0 0 0 x(k) y(k) 1 -x(k)*yprime(k) -yprime(k)*y(k) -yprime(k)];
76         end
77         h = null(bigMatrix);
78         h = h(:,1);
79         H(:,:,i) = reshape(h,3,3);
80
81         Pts_coverH = (Pts_cover' * H(:,:,i))'; %Transform feature points to img
homography
82         dx = Pts_coverH(1,:)./ Pts_coverH(3,:)-Pts_img(1,:);
83         dy = Pts_coverH(2,:)./ Pts_coverH(3,:)-Pts_img(2,:);
84         score(i) = sum(dx.^2+dy.^2 <= thres^2);
85     end
86     [best, bestidx] = max(score);
87     H = H(:,:,bestidx);

```

```

88     cornerArray = [0, 0, 1; size(covers{k},2),0, 1;
size(covers{k},2),size(covers{n},1),1; 0, size(covers{n},1),1]; %corner locations in
homog coords
89     cornerArrayH = (cornerArray * H);
90     cornerArrayH = cornerArrayH ./ cornerArrayH(:,3);
91     cover{n} = cornerArrayH(:,1:2);
92     cover{n} = [cornerArrayH(1,1:2), cornerArrayH(2,1:2), cornerArrayH(3,1:2),
cornerArrayH(4,1:2) ];
93     xyTB(n,:) = [ mean([min(cornerArrayH(:,1)),
max(cornerArrayH(:,1))]),mean([min(cornerArrayH(:,2)), max(cornerArrayH(:,2))]) ] ;
94     end
95     imgAnnotated = img;
96
97     for i = 1:length(cover)
98         imgAnnotated = insertShape(imgAnnotated,'FilledPolygon',cover{i},'Opacity',0.5);
99         imgAnnotated =
insertText(imgAnnotated,xyTB(i,:),names{i},'FontSize',24,'AnchorPoint','Center','BoxC
olor','red');
100     end
101
102     if(display)
103         figure
104         imshow(imgAnnotated);
105     end
106     fileNameEx = char(strcat(num2str(nn,'prob5_classified/%03.f'),".jpg"));
107     imwrite(imgAnnotated,fileNameEx)
108     end

```

Problem 6: Image Stitching

First, the MATLAB tutorial images as stitched together:



And here is my friend doing a pretty nice ski jump! Overall the panorama is a great quality I would say, which I think is because the trees provide a lot of great features to use for matching.



```
1 display = 0; %Show debug images
2
3 % Load images into an array of images imArray
4 buildingImageFiles = fullfile('img_prob6','*.JPG');
5 listImages = dir(buildingImageFiles);
6
7 for i=1:length(listImages)
8     imArray{i} = imread(fullfile(listImages(i).folder, listImages(i).name));
9     imArray{i} = imresize(imArray{i},[1000 NaN]);
10 end
11
12 if(display)
13     montage(imArray)
14 end
15
16 %Total number of images
17 numImages = length(imArray);
18
19 % Read the first image from the image set.
20 I = imArray{1};
21
22 % Find features for I(1) - our "master image"
```

```

23 grayImage = single(rgb2gray(I));
24 [f, d] = vl_sift(grayImage);
25
26 % Initialize variable to hold image sizes.
27 imageSize = zeros(numImages,2);
28 imageSize(1,:) = size(grayImage);
29
30 arrayH = zeros(3,3,numImages); %This is the array of transformations to get the nth
    image on the same homography as the first image. Ie, each layer is the multiplication
    of each previous image
31 arrayH(:,:,1) = eye(3,3); %Make the first transformation matrix the identity
    matrix (the original image doesn't change)
32
33 % Iterate over remaining image pairs
34 for n = 2:numImages
35
36     % Store points and features for I(n-1).
37     fPrevious = f;
38     dPrevious = d;
39
40     % Read I(n).
41     I = imArray{n};
42
43     % Convert image to grayscale.
44     grayImage = single(rgb2gray(I));
45
46     % Save image size.
47     imageSize(n,:) = size(grayImage);
48
49     % Detect features for I(n).
50     [f, d] = vl_sift(grayImage);
51
52     %Find matches between I(n) and I(n-1).
53     [matches, scores] = vl_ubcmatch(dPrevious, d);
54     numMatches = length(scores);
55
56     %Get matching x and y values
57     ptsPrevious = [fPrevious(1:2,matches(1,:));ones(1,numMatches)]; % add ones at
    the end to get homogeneous coordinates
58     ptsI = [f(1:2,matches(2,:));ones(1,numMatches)];
59
60     %Now we estimate the homography
61     n_tries = 50000;
62     thres = 20; %pixels: distance to be considered outlier
63
64     score = zeros(n_tries,1);
65     H = zeros(3,3,n_tries);
66     for i = 1:n_tries
67         x = zeros(1,4);
68         y = zeros(1,4);
69         xprime = zeros(1,4);
70         yprime = zeros(1,4);
71         %Find 4 points at random to estimate homography

```

```

72     for j = 1:4
73         index = randi(numMatches);
74         x(j) = ptsI(1,index);
75         y(j) = ptsI(2,index);
76         xprime(j) = ptsPrevious(1,index);
77         yprime(j) = ptsPrevious(2,index);
78     end
79
80     %Find homography matrix for these 4 points
81     bigMatrix=zeros(8,9);
82     for j=1:2:7
83         k = (j+1)/2;
84         bigMatrix(j:j+1,:) = [x(k) y(k) 1 0 0 0 -x(k)*xprime(k) -xprime(k)*y(k) -
xprime(k);...
85             0 0 0 x(k) y(k) 1 -x(k)*yprime(k) -yprime(k)*y(k) -yprime(k)];
86     end
87     h = null(bigMatrix);
88     h = h(:,1);
89     H(:, :, i) = reshape(h,3,3);
90
91     ptsIH = (ptsI' * H(:, :, i))'; %Transform feature points to img homography
92     dx = ptsIH(1,:)/ ptsIH(3,:)-ptsPrevious(1,:);
93     dy = ptsIH(2,:)/ ptsIH(3,:)-ptsPrevious(2,:);
94     okArray{i} = dx.^2+dy.^2 <= thres^2;
95     score(i) = sum(dx.^2+dy.^2 <= thres^2);
96 end
97 [best, bestidx] = max(score);
98 H = H(:, :, bestidx);
99 arrayH(:, :, n) = H * arrayH(:, :, n-1);
100
101 % -----
102 %                                     Show matches
103 % -----
104 if(display)
105     dh1 = max(size(I,1)-size(imArray{n-1},1),0) ;
106     dh2 = max(size(imArray{n-1},1)-size(I,1),0) ;
107
108     ok = okArray{bestidx};
109     figure
110     subplot(2,1,1) ;
111     imagesc([padarray(imArray{n-1},dh1,'post') padarray(I,dh2,'post')]) ;
112     o = size(imArray{n-1},2) ;
113     line([fPrevious(1,matches(1,:));f(1,matches(2,:))+o], ...
114         [fPrevious(2,matches(1,:));f(2,matches(2,:))]) ;
115     title(sprintf('%d tentative matches', numMatches)) ;
116     axis image off ;
117
118     subplot(2,1,2) ;
119     imagesc([padarray(imArray{n-1},dh1,'post') padarray(I,dh2,'post')]) ;
120     o = size(imArray{n-1},2) ;
121     line([fPrevious(1,matches(1,ok));f(1,matches(2,ok))+o], ...
122         [fPrevious(2,matches(1,ok));f(2,matches(2,ok))]) ;
123     title(sprintf('%d (%.2f%%) inliner matches out of %d', ...

```

```

124         sum(ok), ...
125         100*sum(ok)/numMatches, ...
126         numMatches)) ;
127     axis image off ;
128     drawnow ;
129 end
130 end
131
132 %Now we wish to re-transform our panorama such that the center image is the
133 %"normal" one. This minimizes distortion in edge cases.
134
135 %Compute the output limits for each transform
136 for i = 1:numImages
137     [xlim(i,:), ylim(i,:)] = outputLimits(projective2d(arrayH(:,:,i)), [1
138     imageSize(i,2)], [1 imageSize(i,1)]);
139 end
140 %find the average x limits of each picture. this is essentially the
141 %centroid of each transformed picture. Assuming these are side by side
142 %pictures
143
144 avgXLim = mean(xlim,2);
145 [~, idx] = sort(avgXLim);
146 centerIdx = floor((numImages+1)/2);
147 centerImageIdx = idx(centerIdx);
148
149 Hinv = inv(arrayH(:,:,centerImageIdx));
150
151 for i = 1:numImages
152     arrayH(:,:,i) = arrayH(:,:,i) * Hinv;
153 end
154
155
156 %Stitch our images together for the panorama
157
158 %Find the maximum output limits for the panorma to compute size
159 for i = 1:numImages
160     [xlim(i,:), ylim(i,:)] = outputLimits(projective2d(arrayH(:,:,i)), [1
161     imageSize(i,2)], [1 imageSize(i,1)]);
162 end
163
164 maxImageSize = max(imageSize);
165
166 % Find the minimum and maximum output limits
167 xMin = min([1; xlim(:)]);
168 xMax = max([maxImageSize(2); xlim(:)]);
169
170 yMin = min([1; ylim(:)]);
171 yMax = max([maxImageSize(1); ylim(:)]);
172
173 % width and height of panorama.
174 width = round(xMax - xMin);
175 height = round(yMax - yMin);

```

```

175
176 % Initialize the "empty" panorama.
177 panorama = zeros([height width 3], 'like', I);
178
179 %Now we add our pictures!
180
181 blender = vision.AlphaBlender('Operation', 'Binary mask', ...
182     'MaskSource', 'Input port');
183
184 % Create a 2-D spatial reference object defining the size of the panorama.
185 xLimits = [xMin xMax];
186 yLimits = [yMin yMax];
187 panoramaView = imref2d([height width], xLimits, yLimits);
188
189 % Create the panorama.
190 for i = [1 2 4 5 3]
191
192     I = imArray{i};
193
194     % Transform I into the panorama.
195     warpedImage = imwarp(I, projective2d(arrayH(:,:,i)), 'OutputView', panoramaView);
196
197     %Generate a binary mask.
198     mask = imwarp(true(size(I,1),size(I,2)), projective2d(arrayH(:,:,i)),
199         'OutputView', panoramaView);
200
201     %Overlay the warpedImage onto the panorama.
202     panorama = step(blender, panorama, warpedImage, mask);
203     if(0)
204         figure, imshow(warpedImage)
205     end
206 end
207
208 figure
209 imshow(panorama)

```