

Introduction of Machine Learning

Chul Min Yeum

Assistant Professor

Civil and Environmental Engineering

University of Waterloo, Canada

CIVE 497 – CIVE 700: Smart Structure Technology



UNIVERSITY OF WATERLOO
FACULTY OF ENGINEERING

Last updated: 2019-03-25

What is Machine Learning?

Learning is any process by which a system improves performance from experience (Herbert Simon)

Definition by Tom Mitchell (1998):

model, coefficient

Machine Learning is the study of algorithm that

- improve their performance P
- at some task T
- with experience E

A well-defined learning task is given by $\langle P, T, E \rangle$

Defining the Learning Task

Improve on task T, with respect to performance metric P, based on experience E

T: Playing checkers

P: Percentage of games won against an arbitrary opponent

E: Playing practice games against itself

T: Recognizing hand-written words

P: Percentage of words correctly classified

E: Database of human-labeled images of handwritten words

T: Driving on four-lane highways using vision sensors

P: Average distance traveled before a human-judged error

E: A sequence of images and steering commands recorded while observing a human driver.

T: Categorize email messages as spam or legitimate.

P: Percentage of email messages correctly classified.


E: Database of emails, some with human-given labels

Why is Machine Learning Different from Traditional Programming?

Traditional Programming



Machine Learning




Example: H Beam Classification

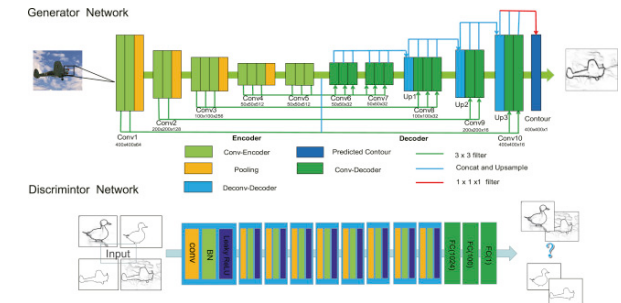

Traditional Programming



Edge detection, Hough transform



Machine Learning



Example: Image-to-Image Translation with Conditional Adversarial Nets



<https://phillipi.github.io/pix2pix/>

Why do We Use Machine Learning?

ML is used when:

- Human expertise does not exist (navigating on Mars)
- Humans can't explain their expertise (speech recognition, image recognition)
- Models are based on huge amounts of data (genomics), which have complex patterns.

- Web search
- Computational biology
- Finance
- E-commerce
- Space exploration
- Robotics
- Information extraction
- Social network
- Debugging Software
- Inspection


What are your applications?

Types of Learning

- **Supervised (inductive) learning**
 - Given: training data + desired outputs (labels)
- **Unsupervised learning**
 - Given: training data (without desired outputs)
- **Semi-supervised learning**
 - Given: training data + a few desired outputs
- **Reinforcement learning**
 - Rewards from sequence of actions

Supervised Learning: Regression

- Given $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$
- Learn a function $f(x)$ to predict y given x
 - y is real-valued == regression

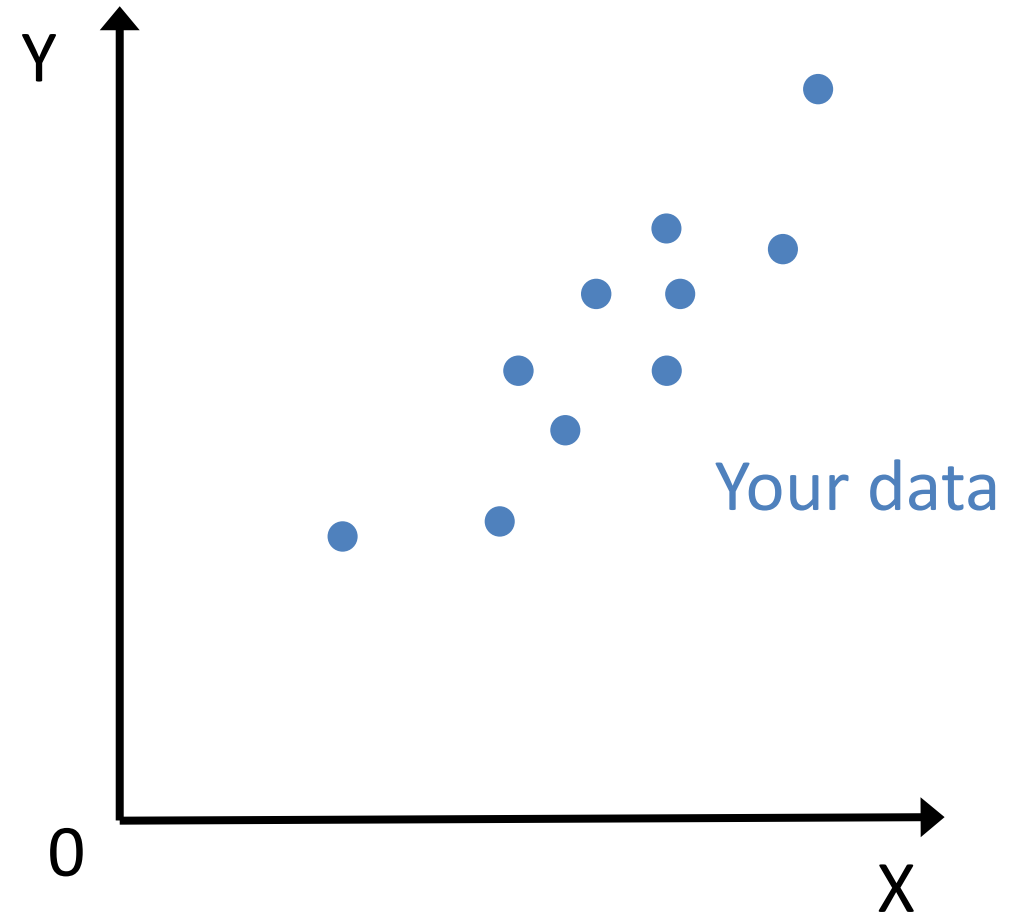


Revisit: Line Fitting

Data (measurement): $(x_1, y_1), \dots, (x_n, y_n)$


Known model: Line $(y_i = m x_i + b)$

We will find m and b .




Supervised Learning: Classification

- Given $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$
- Learn a function $f(x)$ to predict y given x
 - y is categorical == classification





Supervised Learning: Classification

- x can be multi-dimensional
 - Each dimension corresponds to an attribute




- Clump Thickness
- Uniformity of Cell Size
- Uniformity of Cell Shape
- ...



Unsupervised Learning: Clustering


- Given x_1, x_2, \dots, x_n (without labels)
- Output hidden structure behind the x 's
 - E.g., clustering




Example: Unsupervised Learning



Market segmentation




Social network analysis




Autoencoder

- Given a sequence of states and actions with (delayed) rewards, output a policy
 - Policy is a mapping from states \rightarrow actions that tells you what to do in a given state
- Examples:
 - Credit assignment problem
 - Game playing
 - Robot in a maze
 - Balance a pole on your hand



Example: Regression (Supervised Learning)



- Suppose we are given a training set of N observations

$$(x_1, \dots, x_N) \text{ and } (y_1, \dots, y_N), x_i, y_i \in \mathbb{R}$$

- Regression problem is to estimate $y(x)$ from this data

Polynomial Curve Fitting


- The green curve is the true function (which is not a polynomial)
- The data points are uniform in x but have noise in y .
- We will use a **loss function** that measures the squared error in the prediction of $y(x)$ from x . The loss for the red polynomial is the sum of the squared vertical errors.

$$E(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^N \{y(x_i, \mathbf{w}) - t_i\}^2$$

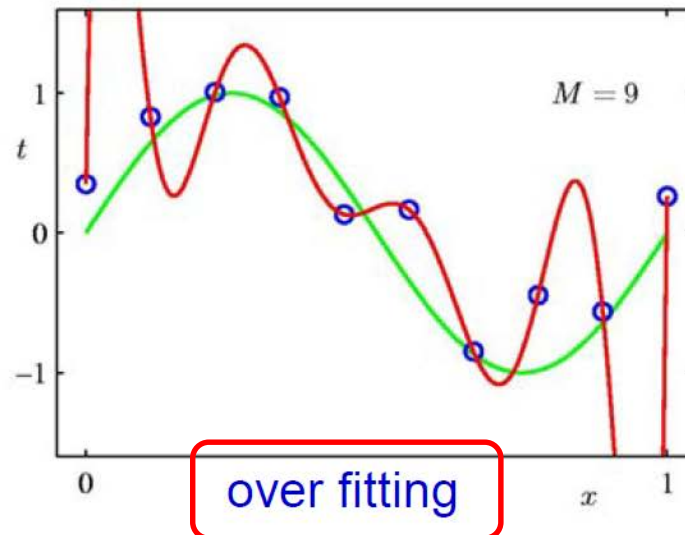



↑
target value

$$y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j$$

polynomial regression



Some Fits to the Data: Which is Best?




Polynomial Coefficients

	$M = 0$	$M = 1$	$M = 3$	$M = 9$
w_0^*	0.19	0.82	0.31	0.35
w_1^*		-1.27	7.99	232.37
w_2^*			-25.43	-5321.83
w_3^*			17.37	48568.31
w_4^*				-231639.30
w_5^*				640042.26
w_6^*				-1061800.52
w_7^*				1042400.18
w_8^*				-557682.99
w_9^*				125201.43

Overfitting

- test data: a different sample from the same true function



Root-Mean-Square (RMS) Error: $E_{\text{RMS}} = \sqrt{2E(\mathbf{w}^*)/N}$



- training error goes to zero, but test error increases with M

Trading off Goodness of Fit against Model Complexity

- If the model has as many degrees of freedom as the data, it can fit the training data perfectly
- But the objective in ML is generalization
- Can expect a model to generalize well if it explains the training data surprisingly well given the complexity of the model.



How to Prevent Over-fitting?

- Add more data than the model “complexity”
- For 9th order polynomial:




- Regularization: penalize large coefficient values



$$\tilde{E}(\mathbf{w}) = \underbrace{\frac{1}{2} \sum_{i=1}^N \{y(x_i, \mathbf{w}) - t_i\}^2}_{\text{loss function}} + \underbrace{\frac{\lambda}{2} \|\mathbf{w}\|^2}_{\text{regularization}} \quad \text{“ridge” regression}$$



Learning Rate



$$x_{n+1} = x_n - \frac{f(x_n)}{f'(x_n)}$$



Use a validation set:

Divide the total dataset into three subsets:

- **Training data** is used for learning the parameters of the model.
- **Validation data** is not used for learning but is used for deciding what type of model and what amount of regularization works best.
- **Test data** is used to get a final, unbiased estimate of how well the learning machine works. We expect this estimate to be worse than on the validation data.

We could then re-divide the total dataset to get another unbiased estimate of the true error rate.

Example: Damage Detection




Collapse classification






Pipeline inspection

- Supervised
- Unsupervised
- Semi-supervised



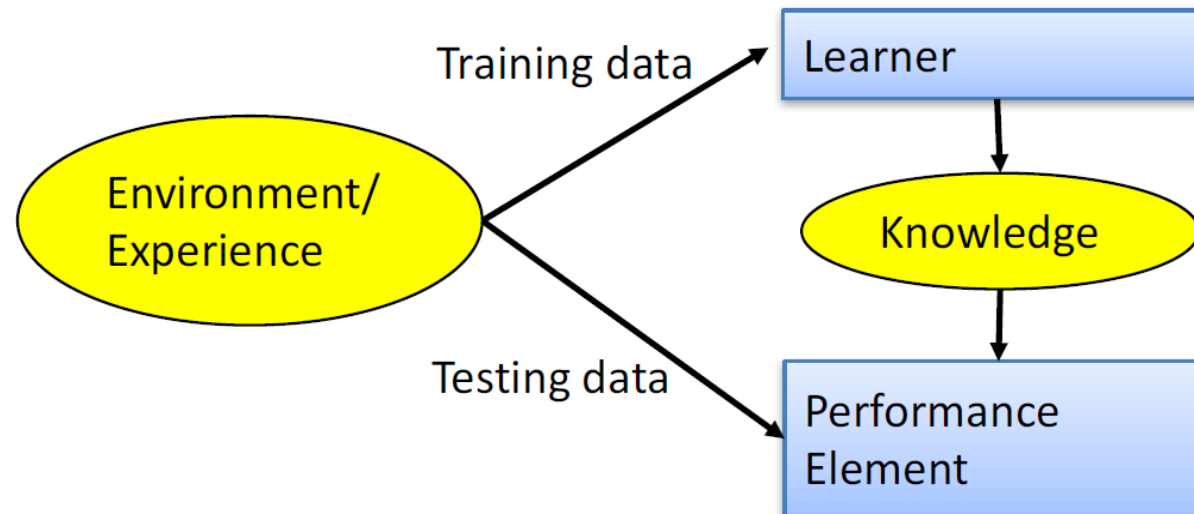
Leak detection (Cody et al, 2018)


Example: Homography Estimation and Fundamental Matrix



Designing a Learning System

- Choose the training experience
- Choose exactly what is to be learned
 - i.e. the **target function**
- Choose how to represent the target function
- Choose a learning algorithm to infer the target function from the experience







Loop

- Understand domain, prior knowledge, and goals
- Data integration, selection, cleaning, pre-processing, etc.
- Learn models
- Interpret results
- Consolidate and deploy discovered knowledge

Curse of Dimensionality



- "As the number of features or dimensions grows, the amount of data we need to generalize accurately grows exponentially," Charles Isbell
- Think of image recognition problem of high resolution images $1280 \times 720 = 921,600$ pixels i.e. 921600 dimensions.
- That's why it's called **Curse of Dimensionality**. Value added by additional dimension is much smaller compared to overhead it adds to the algorithm.

SIFT Descriptor

Example: Summation of N Numbers from a to a+N

$$\sum_{i=1}^N i = \frac{N(N+1)}{2}$$

$$\sum_{i=a}^{N+a} i = \frac{N(a+N+1)}{2}$$

Given $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$

How to design your problem?



Slide Credits and References

- Lecture notes: Eric Eaton
- Lecture notes: A. Zisserman
- Lecture notes: David Sontag
- Bishop (2006) Pattern Recognition and Machine Learning